

```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from scipy import stats
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
```

EDA and Data Processing

```
In [ ]: df = pd.read_csv("dataset/Jamboree_Admission.csv")
df.head()
```

|   | Serial No. | GRE Score | TOEFL Score | University Rating | SOP | LOR | CGPA | Research | Chance of Admit |
|---|------------|-----------|-------------|-------------------|-----|-----|------|----------|-----------------|
| 0 | 1          | 337       | 118         | 4                 | 4.5 | 4.5 | 9.65 | 1        | 0.92            |
| 1 | 2          | 324       | 107         | 4                 | 4.0 | 4.5 | 8.87 | 1        | 0.76            |
| 2 | 3          | 316       | 104         | 3                 | 3.0 | 3.5 | 8.00 | 1        | 0.72            |
| 3 | 4          | 322       | 110         | 3                 | 3.5 | 2.5 | 8.67 | 1        | 0.80            |
| 4 | 5          | 314       | 103         | 2                 | 2.0 | 3.0 | 8.21 | 0        | 0.65            |

```
In [ ]: df.shape
```

```
Out[ ]: (500, 9)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Serial No.          500 non-null    int64
1   GRE Score           500 non-null    int64
2   TOEFL Score         500 non-null    int64
3   University Rating   500 non-null    int64
4   SOP                 500 non-null    float64
5   LOR                 500 non-null    float64
6   CGPA                500 non-null    float64
7   Research            500 non-null    int64
8   Chance of Admit     500 non-null    float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

```
In [ ]: df.describe()
```

|       | Serial No. | GRE Score  | TOEFL Score | University Rating | SOP        | LOR        | CGPA       | Research   | Chance of Admit |
|-------|------------|------------|-------------|-------------------|------------|------------|------------|------------|-----------------|
| count | 500.000000 | 500.000000 | 500.000000  | 500.000000        | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500.000000      |
| mean  | 250.500000 | 316.472000 | 107.192000  | 3.114000          | 3.374000   | 3.48400    | 8.576440   | 0.560000   | 0.72174         |
| std   | 144.481833 | 11.295148  | 6.081868    | 1.143512          | 0.991004   | 0.92545    | 0.604813   | 0.496884   | 0.14114         |
| min   | 1.000000   | 290.000000 | 92.000000   | 1.000000          | 1.000000   | 1.00000    | 6.800000   | 0.000000   | 0.34000         |
| 25%   | 125.750000 | 308.000000 | 103.000000  | 2.000000          | 2.500000   | 3.00000    | 8.127500   | 0.000000   | 0.63000         |
| 50%   | 250.500000 | 317.000000 | 107.000000  | 3.000000          | 3.500000   | 3.50000    | 8.560000   | 1.000000   | 0.72000         |
| 75%   | 375.250000 | 325.000000 | 112.000000  | 4.000000          | 4.000000   | 4.00000    | 9.040000   | 1.000000   | 0.82000         |
| max   | 500.000000 | 340.000000 | 120.000000  | 5.000000          | 5.00000    | 5.00000    | 9.920000   | 1.000000   | 0.97000         |


```
In [ ]: df.drop(["Serial No."], axis=1, inplace=True)
```

```
In [ ]: df.rename(columns=lambda x: x.strip(), inplace=True)
df.columns
```

```
Out[ ]: Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA',
              'Research', 'Chance of Admit'],
              dtype='object')
```

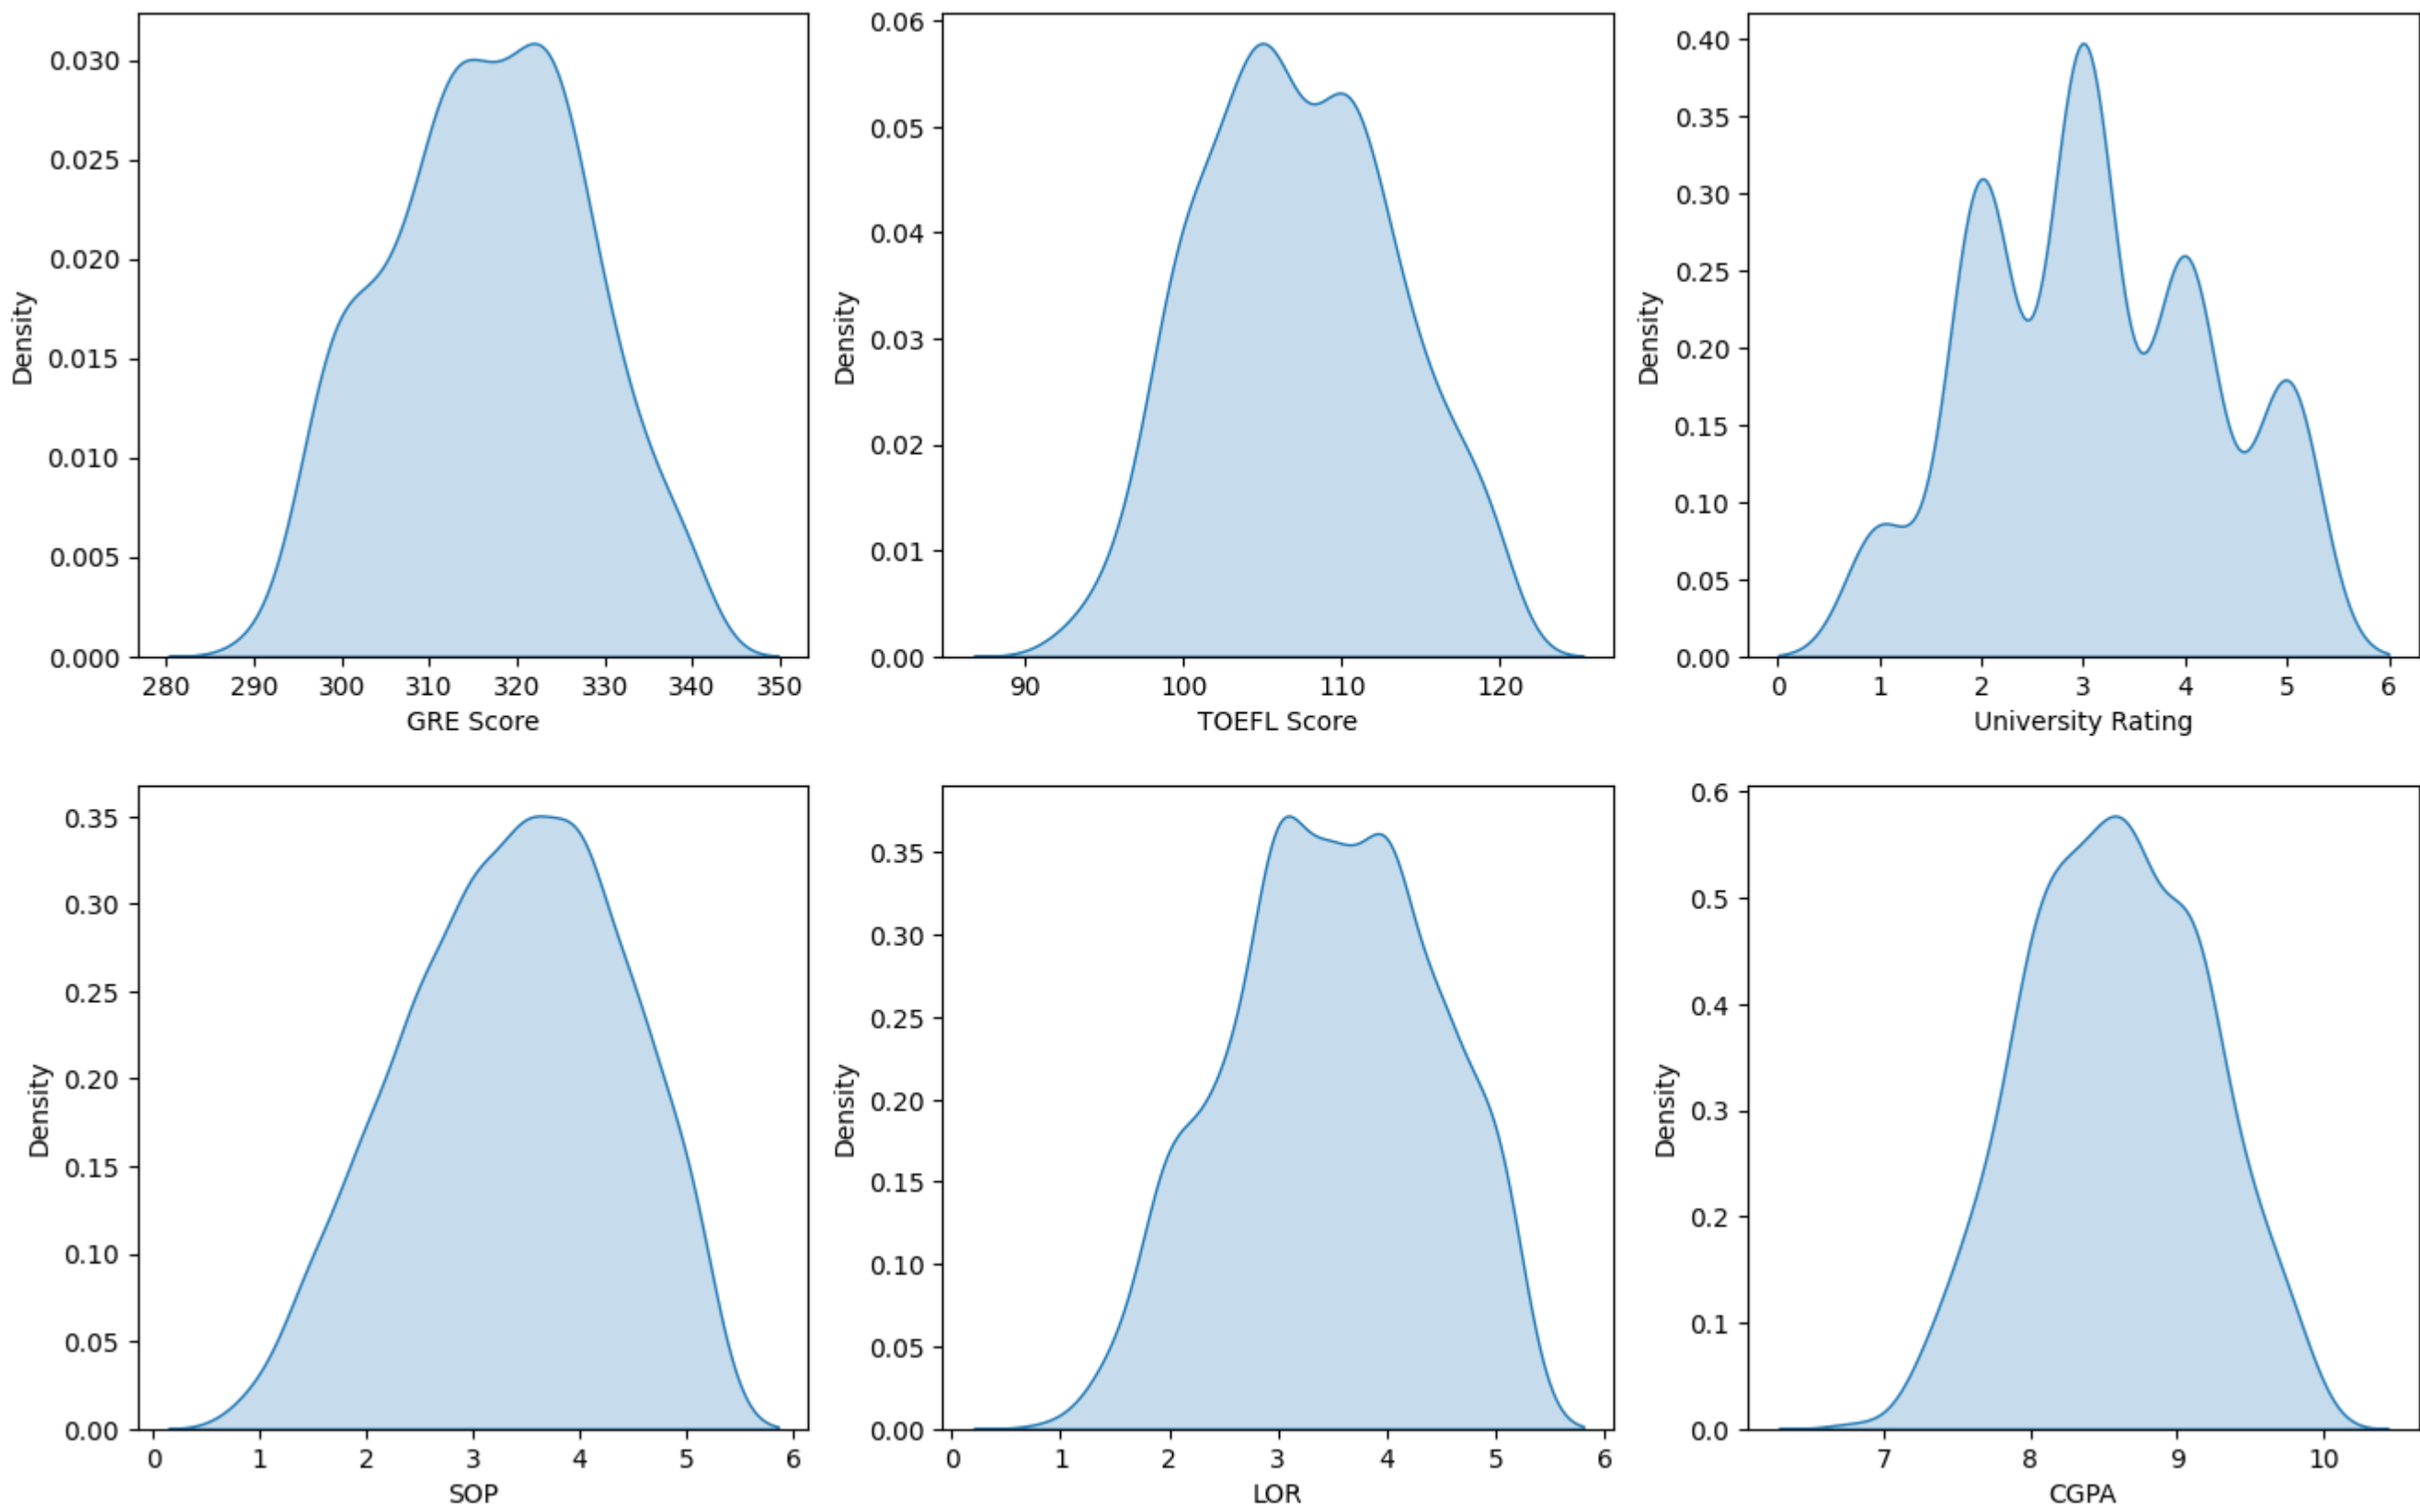
```
In [ ]: sns.kdeplot(df["GRE Score"], fill=True)
```

```
Out[ ]: <Axes: xlabel='GRE Score', ylabel='Density'>
```



```
In [ ]: fig=plt.figure(figsize=(16,10))
```

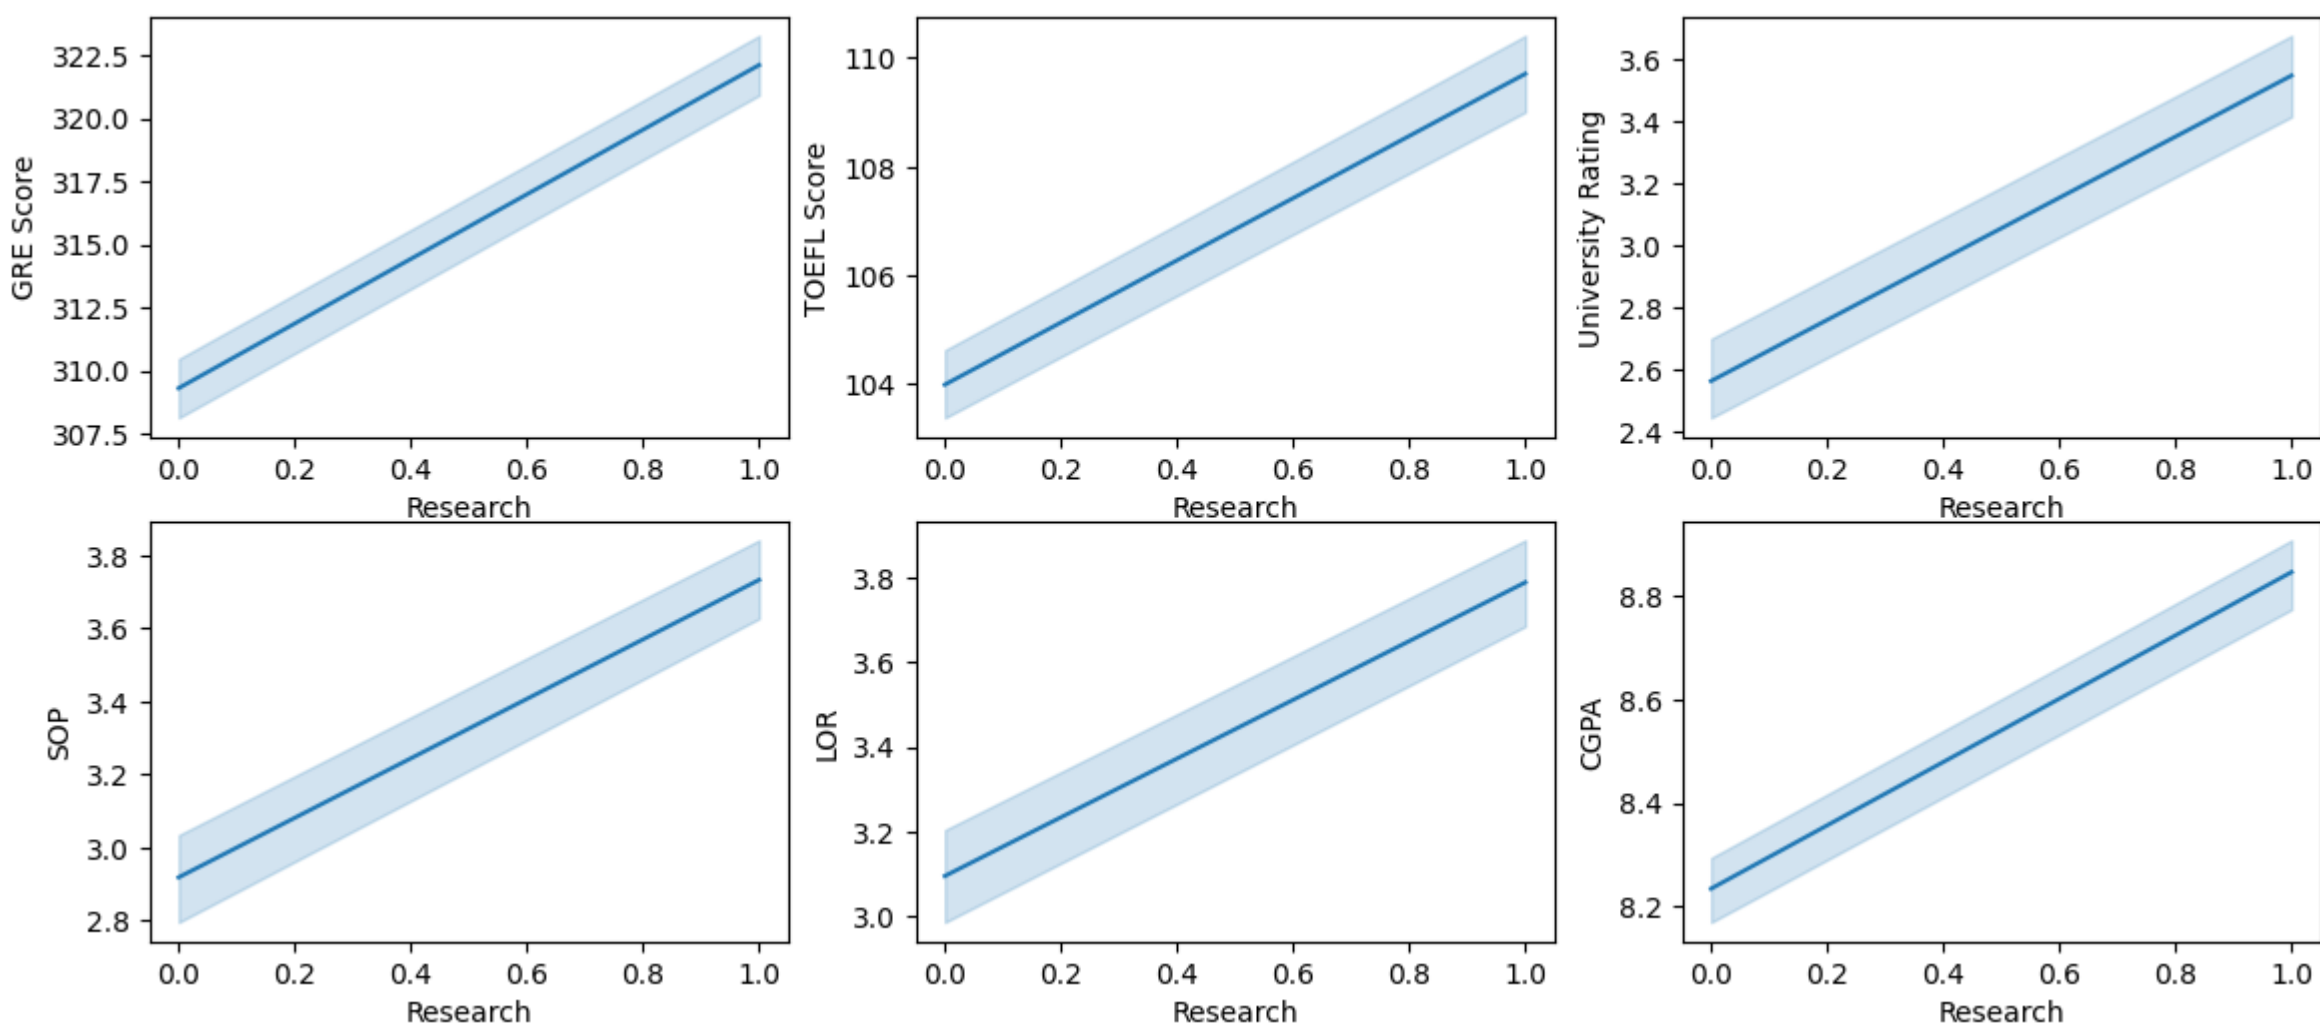
```
for ind_number, col_name in enumerate(df.columns):
    if ind_number<6:
        plt.subplot(2,3,ind_number+1)
        sns.kdeplot(df[col_name], fill=True)
```



```
In [ ]: # Bi Variate Analysis
fig=plt.figure(figsize=(14,6))

for ind_number, col_name in enumerate(df.columns):
    if ind_number<6:
        plt.subplot(2,3,ind_number+1)
        sns.lineplot(x='Research', y=col_name, data=df)

# plt.subplot(1,2,2)
# sns.scatterplot(x='Research', y='TOEFL Score', data=df)
```



```
In [ ]: fig=plt.figure(figsize=(14,8))
```

```
plt.subplot(2,3,1)
sns.lineplot(y='TOEFL Score', x='GRE Score', data=df)

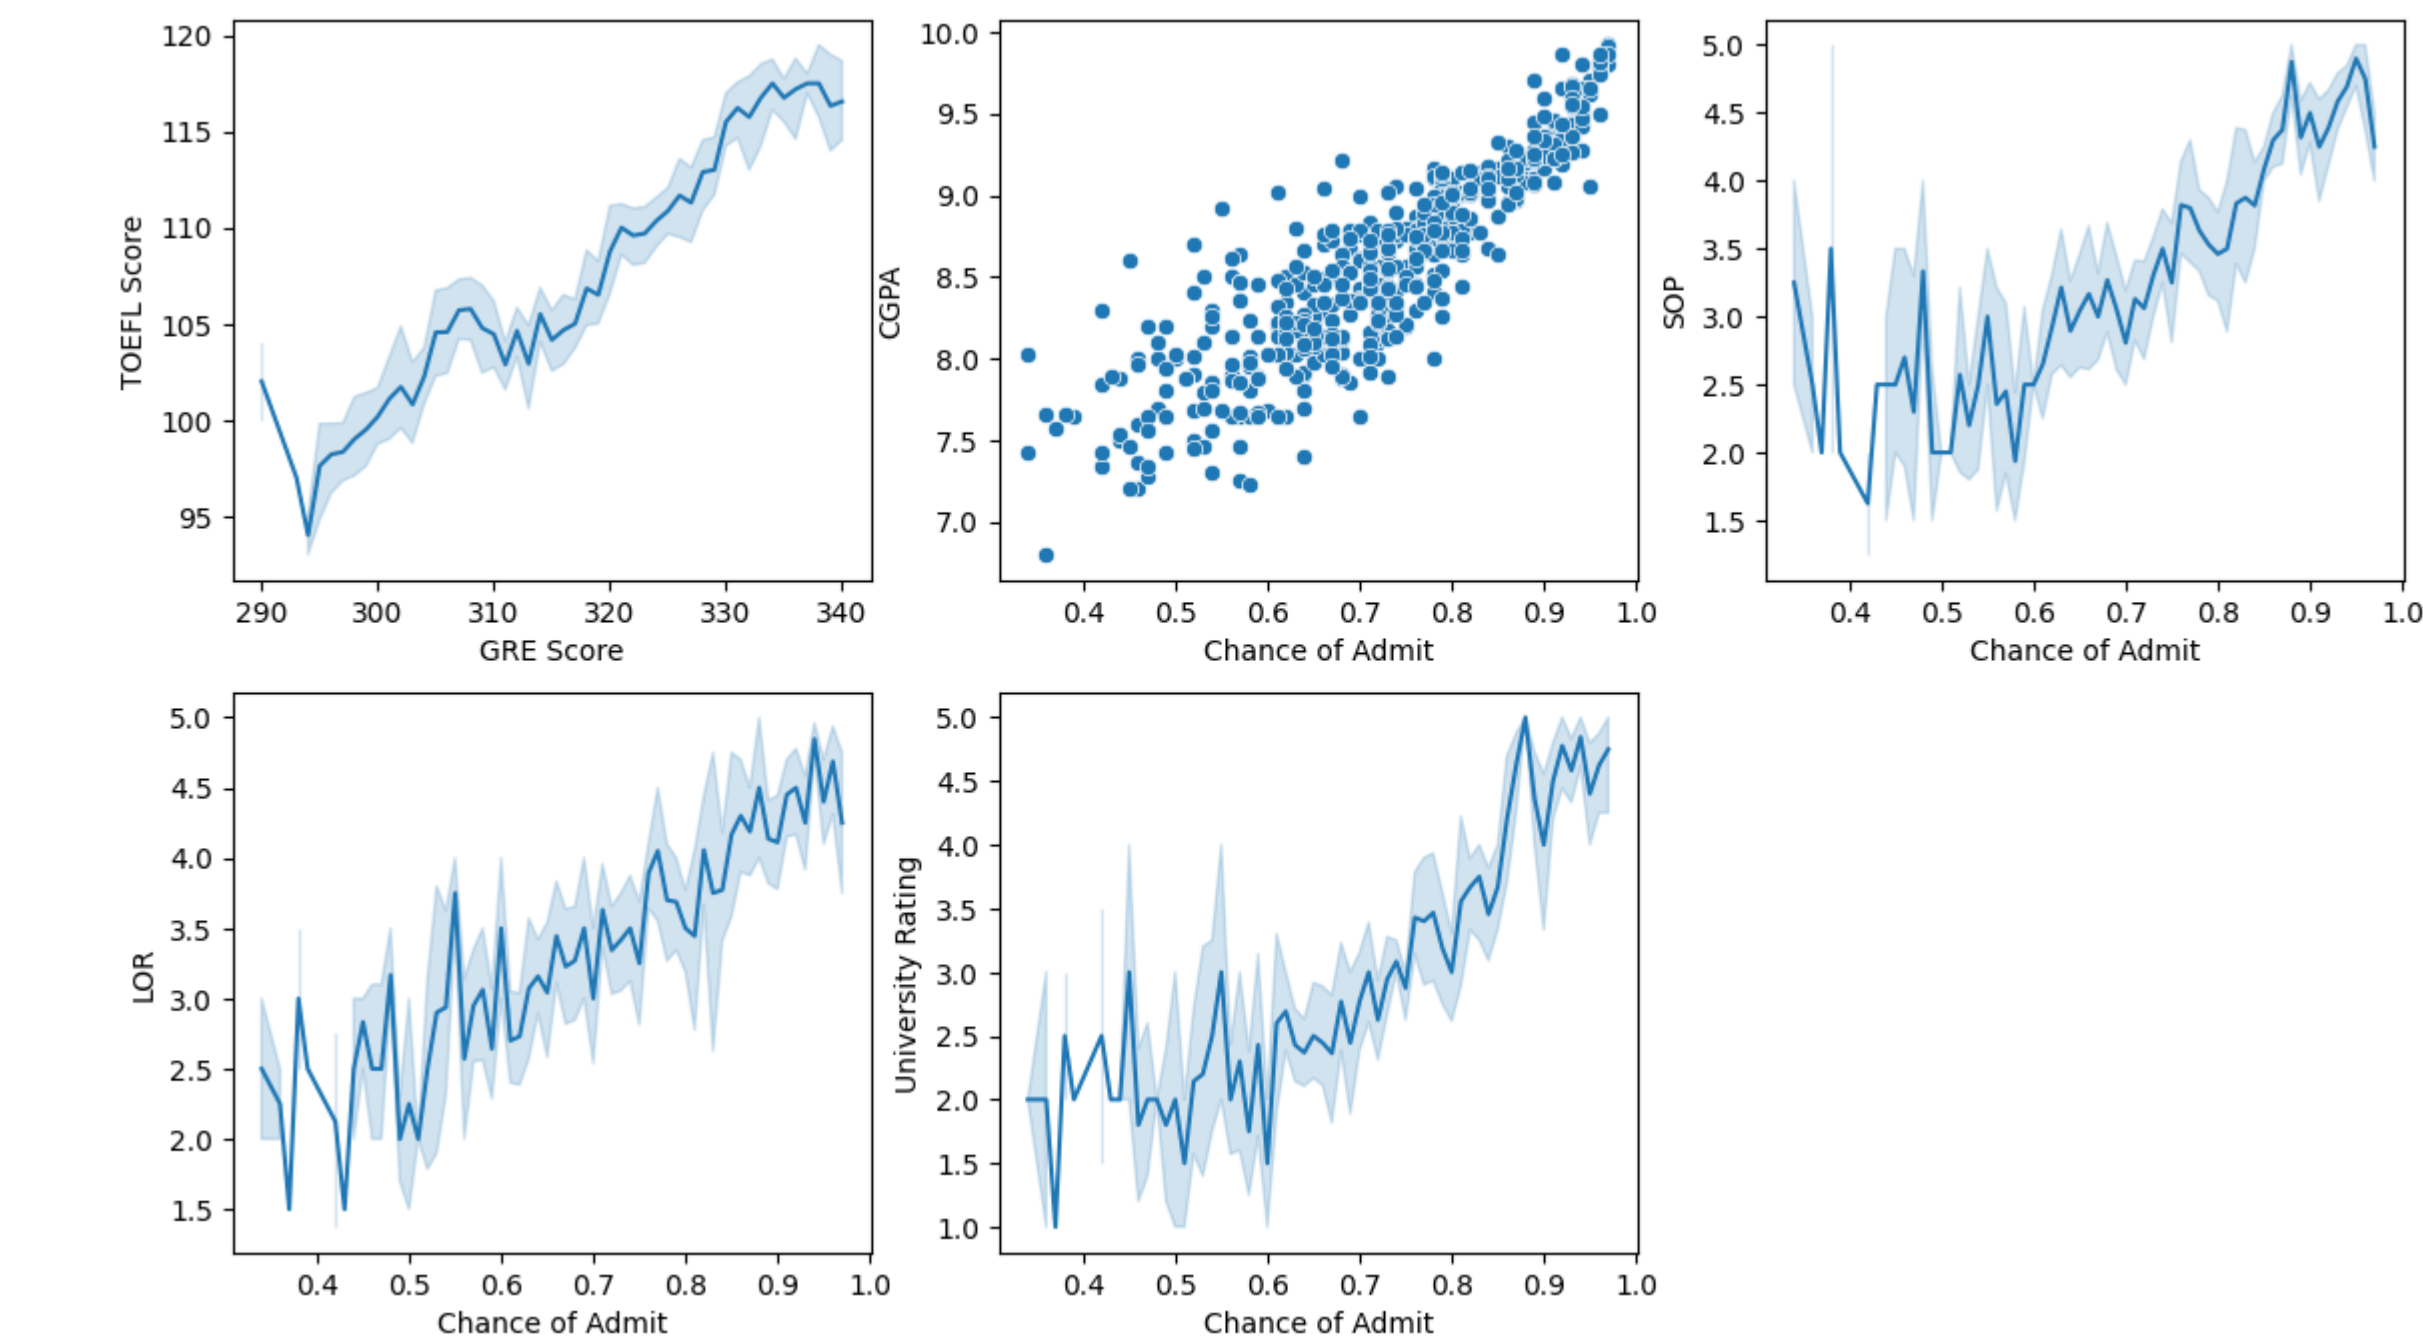
plt.subplot(2,3,2)
sns.scatterplot(y='CGPA', x='Chance of Admit', data=df)

plt.subplot(2,3,3)
sns.lineplot(y='SOP', x='Chance of Admit', data=df)


plt.subplot(2,3,4)
sns.lineplot(y='LOR', x='Chance of Admit', data=df)

plt.subplot(2,3,5)
sns.lineplot(y='University Rating', x='Chance of Admit', data=df)
```

```
Out[ ]: <Axes: xlabel='Chance of Admit', ylabel='University Rating'>
```

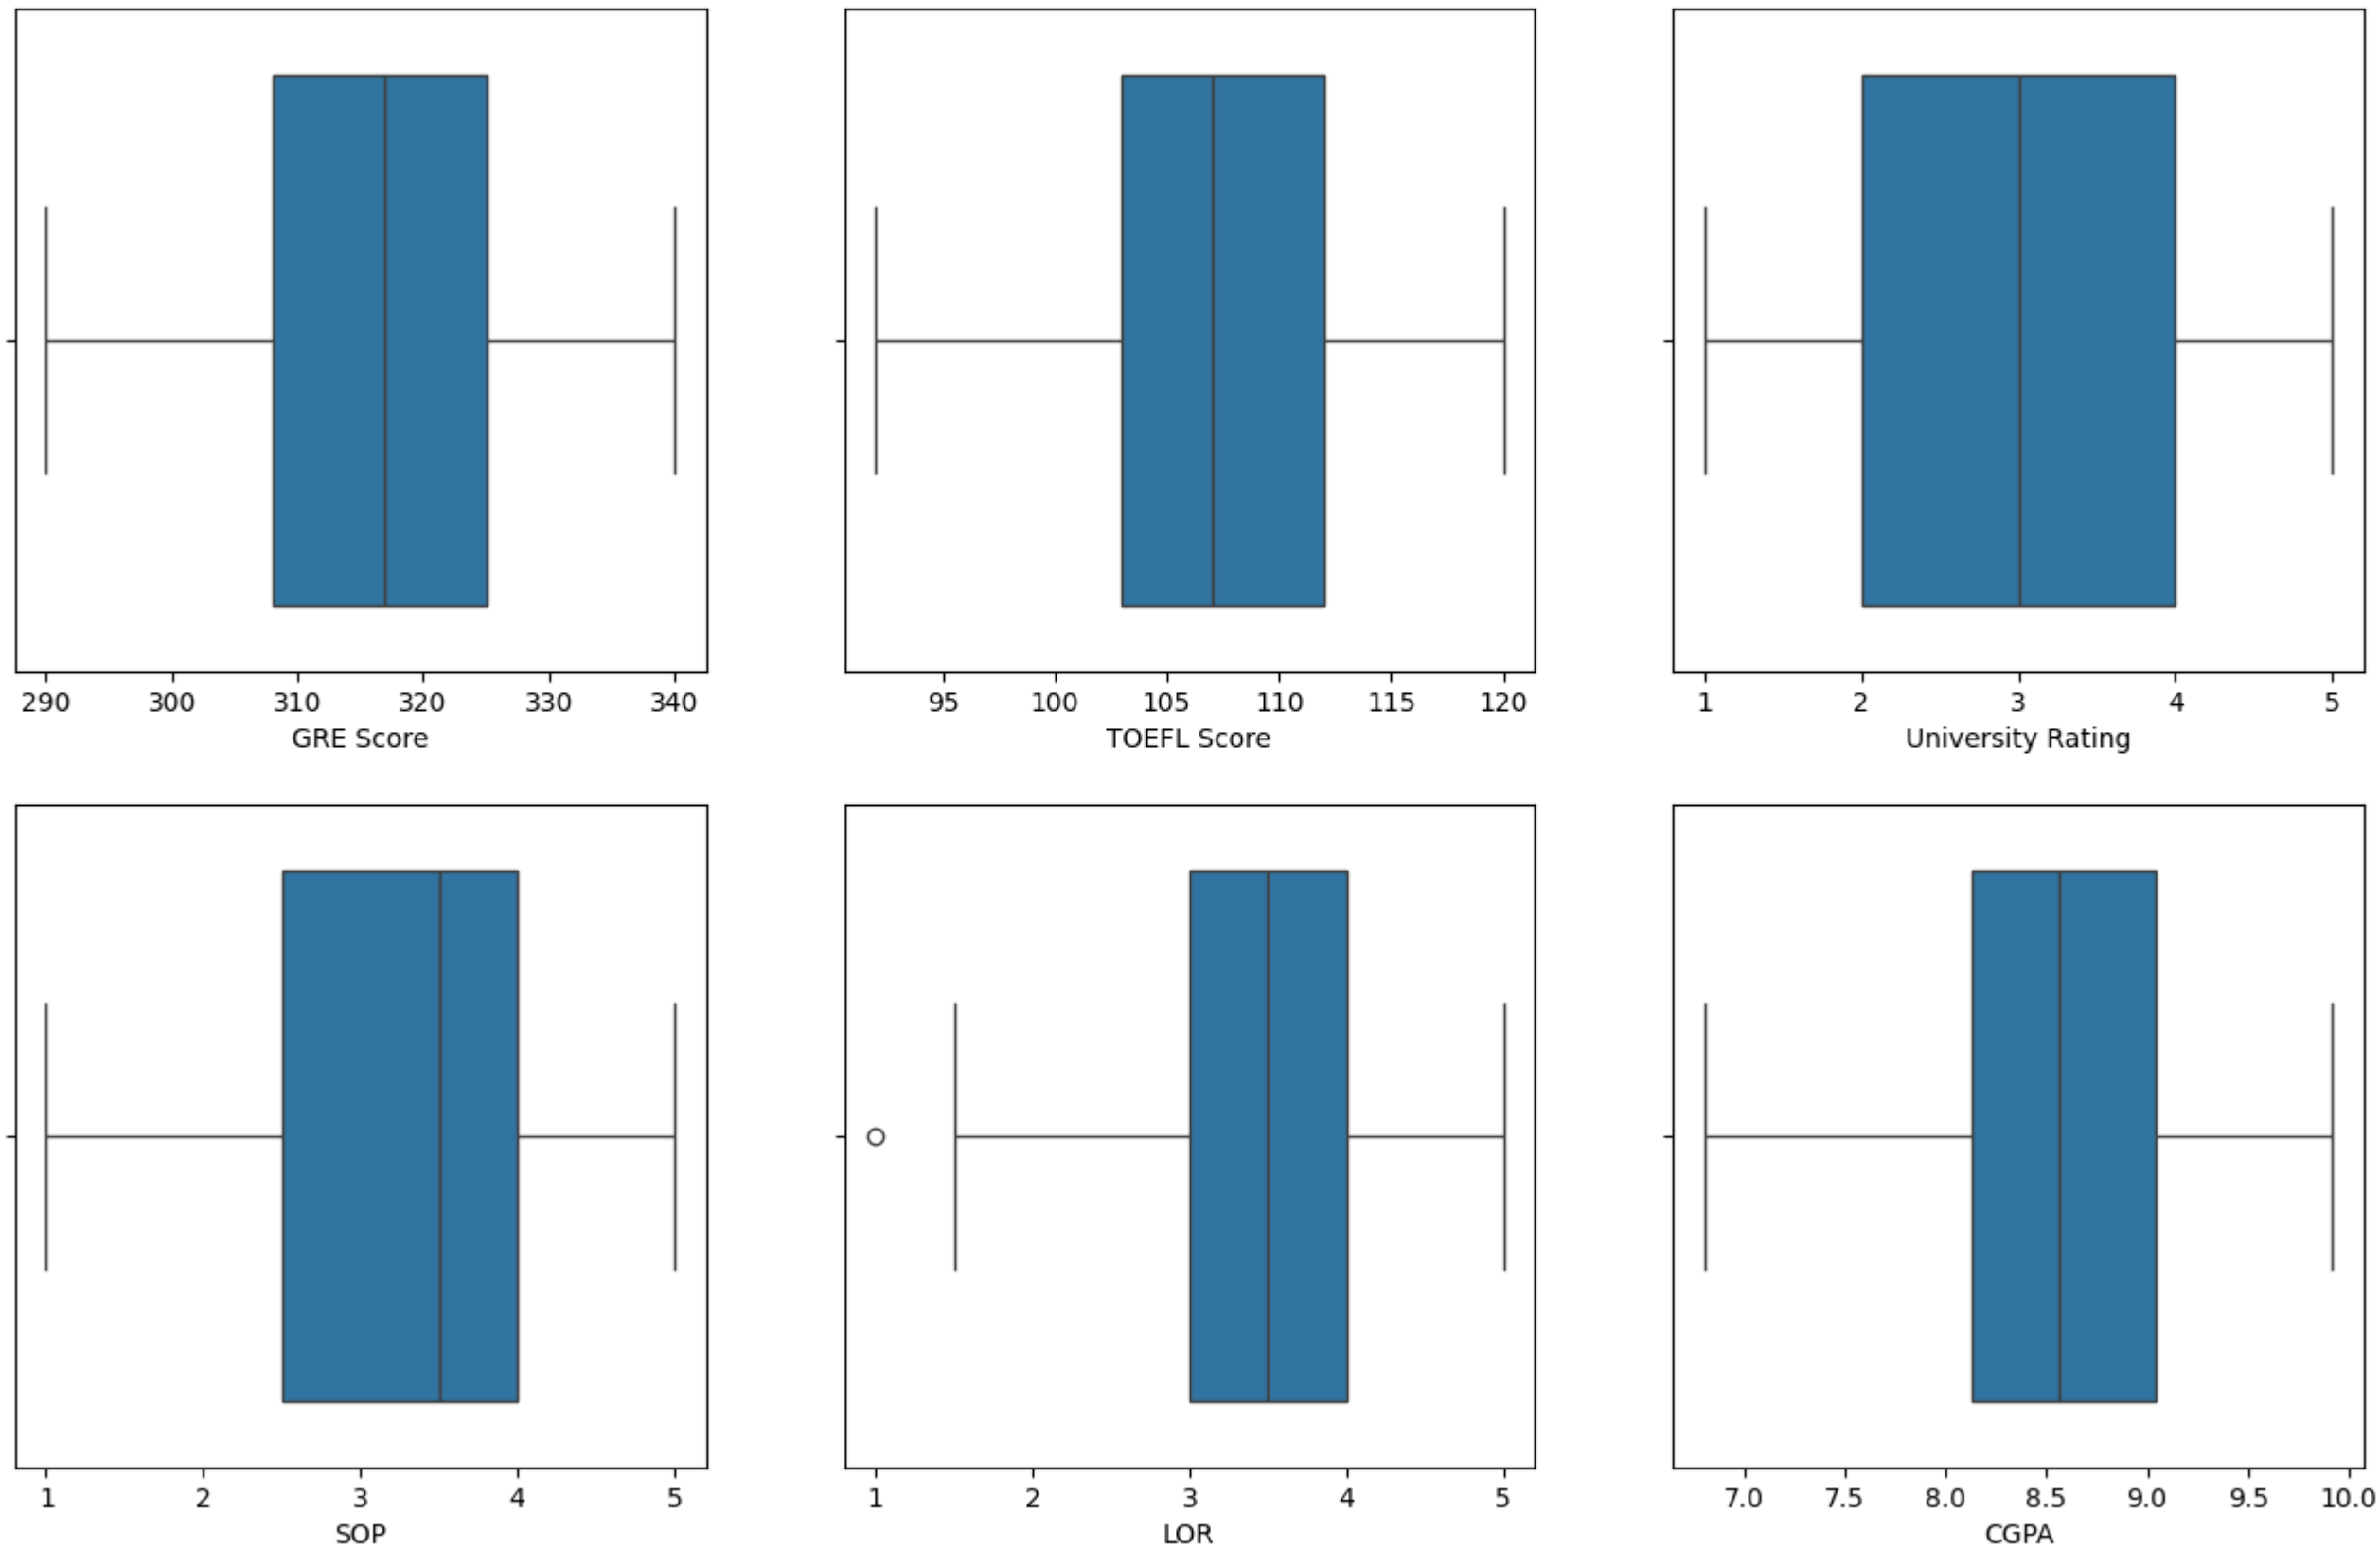


```
In [ ]: sns.heatmap(df.corr(), cmap= "Blues", annot=True)
plt.show()
```



```
In [ ]: fig=plt.figure(figsize=(16,10))
```

```
for ind_number, col_name in enumerate(df.columns):
    if ind_number<6:
        plt.subplot(2,3,ind_number+1)
        sns.boxplot(x=col_name, data=df)
```



```
In [ ]: df.head()
```

|   | GRE Score | TOEFL Score | University Rating | SOP | LOR | CGPA | Research | Chance of Admit |
|---|-----------|-------------|-------------------|-----|-----|------|----------|-----------------|
| 0 | 337       | 118         | 4                 | 4.5 | 4.5 | 9.65 | 1        | 0.92            |
| 1 | 324       | 107         | 4                 | 4.0 | 4.5 | 8.87 | 1        | 0.76            |
| 2 | 316       | 104         | 3                 | 3.0 | 3.5 | 8.00 | 1        | 0.72            |
| 3 | 322       | 110         | 3                 | 3.5 | 2.5 | 8.67 | 1        | 0.80            |
| 4 | 314       | 103         | 2                 | 2.0 | 3.0 | 8.21 | 0        | 0.65            |

## Multicollinearity and VIF

```
In [ ]: df_test = df.copy()

# Standard Scaler
numeric_columns = ['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA', 'Research']
scaler = StandardScaler()
scaler.fit(df_test[numeric_columns])
df_test[numeric_columns] = scaler.transform(df_test[numeric_columns])

## Min-Max Scaler Data Preparation
# scaler = MinMaxScaler()
# df_test = pd.DataFrame(scaler.fit_transform(df_test), columns=df_test.columns)
# df_test.head()
```

```
In [ ]: y = df_test[['Chance of Admit']]
x = df_test.drop(['Chance of Admit'], axis=1)

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1, random_state=2)
x_train
```

|     | GRE Score | TOEFL Score | University Rating | SOP       | LOR       | CGPA      | Research  |
|-----|-----------|-------------|-------------------|-----------|-----------|-----------|-----------|
| 407 | -1.637030 | -1.183716   | -0.099793         | -0.882817 | 0.558125  | -1.036796 | 0.886405  |
| 291 | -1.459786 | -0.854540   | -0.975168         | -1.892906 | -1.605151 | -1.169201 | -1.128152 |
| 259 | 1.287504  | 1.943453    | 0.775582          | 1.642404  | 1.098944  | 1.263738  | 0.886405  |
| 231 | 0.224037  | -0.196189   | -0.099793         | 0.127271  | -1.064332 | -0.407873 | 0.886405  |
| 478 | 0.135415  | -0.689952   | -0.099793         | 0.632315  | 1.098944  | -0.143063 | 0.886405  |
| ... | ...       | ...         | ...               | ...       | ...       | ...       | ...       |
| 22  | 1.021637  | 1.449690    | 1.650957          | 1.642404  | 1.639763  | 1.528547  | 0.886405  |
| 72  | 0.401282  | 0.626751    | 1.650957          | 1.642404  | 1.639763  | 1.445794  | 0.886405  |
| 493 | -1.459786 | -2.006655   | -0.975168         | -0.377773 | -2.145970 | -0.589930 | 0.886405  |
| 15  | -0.219074 | -0.360777   | -0.099793         | 0.127271  | -1.064332 | -0.457525 | -1.128152 |
| 168 | -2.080142 | -1.677479   | -0.975168         | -1.387862 | 0.558125  | -1.285055 | 0.886405  |

450 rows × 7 columns

```
In [ ]: # VIF
vif = pd.DataFrame()
vif['Features'] = x_train.columns
vif['VIF'] = [variance_inflation_factor(x_train.values, i) for i in range(x_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

|   | Features          | VIF  |
|---|-------------------|------|
| 5 | CGPA              | 4.71 |
| 0 | GRE Score         | 4.35 |
| 1 | TOEFL Score       | 3.83 |
| 3 | SOP               | 2.76 |
| 2 | University Rating | 2.60 |
| 4 | LOR               | 2.01 |
| 6 | Research          | 1.46 |

```
In [ ]: vif_thr = 5
r2_thr = 0.85
i=1
features_removed = ['CGPA']

def adj_r2_func(r2, x):
    return 1 - (1-r2) * (x.shape[0]-1)/(x.shape[0]-x.shape[1]-1)

while True:
    vif = pd.DataFrame()
    vif['Features'] = x_train.columns
    vif['VIF'] = [variance_inflation_factor(x_train.values, i) for i in range(x_train.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)

    x_train = x_train[vif['Features'][1:].values]
    x_test = x_test[vif['Features'][1:].values]

    model = LinearRegression()
    model.fit(x_train, y_train)

    adj_r2 = adj_r2_func(r2_score(y_test, model.predict(x_test)), x_test)

    if (vif.iloc[0]['VIF'] < vif_thr) or (adj_r2 < r2_thr):
        print('Reached threshold')
        print('Highest vif:',vif.iloc[0]['Features'])
        print('Current adj_R2:',adj_r2)
        print('Features removed:', 1)
        print('List of features removed:', features_removed)
        break

    features_removed.append(vif.iloc[0]['Features'])
    i+=1
```

Reached threshold  
Highest vif: CGPA  
Current adj\_R2 0.7610620551043735  
Features removed: 1  
List of features removed: ['CGPA']

```
In [ ]:
```

## Model Implementation

Stats Model



In [ ]: df\_stats = df.copy()  
df\_stats.head()

|   | GRE Score | TOEFL Score | University Rating | SOP | LOR | CGPA | Research | Chance of Admit |      |
|---|-----------|-------------|-------------------|-----|-----|------|----------|-----------------|------|
| 0 | 337       | 118         |                   | 4   | 4.5 | 4.5  | 9.65     | 1               | 0.92 |
| 1 | 324       | 107         |                   | 4   | 4.0 | 4.5  | 8.87     | 1               | 0.76 |
| 2 | 316       | 104         |                   | 3   | 3.0 | 3.5  | 8.00     | 1               | 0.72 |
| 3 | 322       | 110         |                   | 3   | 3.5 | 2.5  | 8.67     | 1               | 0.80 |
| 4 | 314       | 103         |                   | 2   | 2.0 | 3.0  | 8.21     | 0               | 0.65 |

In [ ]: # Standard Scaler  
numeric\_columns = ['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA', 'Research']  
scaler = StandardScaler()  
scaler.fit(df\_stats[numeric\_columns])  
df\_stats[numeric\_columns] = scaler.transform(df\_stats[numeric\_columns])  
  
## Min-Max Scaler Data Preparation  
# scaler = MinMaxScaler()  
# df\_stats = pd.DataFrame(scaler.fit\_transform(df\_stats), columns=df\_stats.columns)  
# df\_stats.head()

In [ ]: y = df\_stats[['Chance of Admit']]  
x = df\_stats.drop(['Chance of Admit'], axis=1)  
  
x\_train, x\_test, y\_train, y\_test = train\_test\_split(x, y, test\_size=0.05, random\_state=2)  
x\_train

|     | GRE Score | TOEFL Score | University Rating | SOP       | LOR       | CGPA      | Research  |
|-----|-----------|-------------|-------------------|-----------|-----------|-----------|-----------|
| 171 | 1.553371  | 1.614278    | 1.650957          | 0.632315  | 1.098944  | 0.816871  | 0.886405  |
| 18  | 0.135415  | 0.462163    | -0.099793         | 0.632315  | -0.523513 | 0.370005  | -1.128152 |
| 350 | 0.135415  | -0.031601   | -0.099793         | -0.377773 | 0.017306  | -0.507177 | 0.886405  |
| 476 | -1.105297 | -0.525364   | -0.099793         | -0.882817 | -1.605151 | -0.755436 | -1.128152 |
| 20  | -0.396319 | -0.031601   | -0.099793         | -0.377773 | -1.605151 | -1.119549 | 0.886405  |
| ... | ...       | ...         | ...               | ...       | ...       | ...       | ...       |
| 22  | 1.021637  | 1.449690    | 1.650957          | 1.642404  | 1.639763  | 1.528547  | 0.886405  |
| 72  | 0.401282  | 0.626751    | 1.650957          | 1.642404  | 1.639763  | 1.445794  | 0.886405  |
| 493 | -1.459786 | -2.006655   | -0.975168         | -0.377773 | -2.145970 | -0.589930 | 0.886405  |
| 15  | -0.219074 | -0.360777   | -0.099793         | 0.127271  | -1.064332 | -0.457525 | -1.128152 |
| 168 | -2.080142 | -1.677479   | -0.975168         | -1.387862 | 0.558125  | -1.285055 | 0.886405  |

475 rows × 7 columns

In [ ]: y\_train = np.array(y\_train)  
x\_train = sm.add\_constant(x\_train)  
x\_train

|     | const | GRE Score | TOEFL Score | University Rating | SOP       | LOR       | CGPA      | Research  |
|-----|-------|-----------|-------------|-------------------|-----------|-----------|-----------|-----------|
| 171 | 1.0   | 1.553371  | 1.614278    | 1.650957          | 0.632315  | 1.098944  | 0.816871  | 0.886405  |
| 18  | 1.0   | 0.135415  | 0.462163    | -0.099793         | 0.632315  | -0.523513 | 0.370005  | -1.128152 |
| 350 | 1.0   | 0.135415  | -0.031601   | -0.099793         | -0.377773 | 0.017306  | -0.507177 | 0.886405  |
| 476 | 1.0   | -1.105297 | -0.525364   | -0.099793         | -0.882817 | -1.605151 | -0.755436 | -1.128152 |
| 20  | 1.0   | -0.396319 | -0.031601   | -0.099793         | -0.377773 | -1.605151 | -1.119549 | 0.886405  |
| ... | ...   | ...       | ...         | ...               | ...       | ...       | ...       | ...       |
| 22  | 1.0   | 1.021637  | 1.449690    | 1.650957          | 1.642404  | 1.639763  | 1.528547  | 0.886405  |
| 72  | 1.0   | 0.401282  | 0.626751    | 1.650957          | 1.642404  | 1.639763  | 1.445794  | 0.886405  |
| 493 | 1.0   | -1.459786 | -2.006655   | -0.975168         | -0.377773 | -2.145970 | -0.589930 | 0.886405  |
| 15  | 1.0   | -0.219074 | -0.360777   | -0.099793         | 0.127271  | -1.064332 | -0.457525 | -1.128152 |
| 168 | 1.0   | -2.080142 | -1.677479   | -0.975168         | -1.387862 | 0.558125  | -1.285055 | 0.886405  |

475 rows × 8 columns

In [ ]: x\_train.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Index: 475 entries, 171 to 168  
Data columns (total 8 columns):  
# Column Non-Null Count Dtype  
--- --- -  
0 const 475 non-null float64  
1 GRE Score 475 non-null float64  
2 TOEFL Score 475 non-null float64  
3 University Rating 475 non-null float64  
4 SOP 475 non-null float64  
5 LOR 475 non-null float64  
6 CGPA 475 non-null float64  
7 Research 475 non-null float64  
dtypes: float64(8)  
memory usage: 33.4 KB

In [ ]: model = sm.OLS(y\_train, x\_train)  
results = model.fit()  
  
# Summary statistics of the model  
print(results.summary())

```

OLS Regression Results
=====
Dep. Variable:          y      R-squared:          0.825
Model:                OLS     Adj. R-squared:      0.822
Method:               Least Squares      F-statistic:    313.5
Date:                Thu, 07 Dec 2023      Prob (F-statistic):  5.08e-172
Time:                18:26:03      Log-Likelihood:    673.17
No. Observations:      475      AIC:              -1330.
Df Residuals:          467      BIC:              -1297.
Df Model:              7
Covariance Type:       nonrobust
=====
coef    std err          t      P>|t|      [0.025    0.975]
-----
const      0.7221         0.003    265.958      0.000      0.717      0.727
GRE Score   0.0243         0.006      4.245      0.000      0.013      0.036
TOEFL Score  0.0152         0.005      2.823      0.005      0.005      0.026
University Rating  0.0078         0.004      1.797      0.073      -0.001      0.016
SOP         0.0003         0.005      0.076      0.939     -0.009      0.009
LOR         0.0168         0.004      4.336      0.000      0.009      0.024
CGPA        0.0783         0.006     11.915      0.000      0.059      0.092
Research    0.0108         0.003      3.270      0.001      0.004      0.017
=====
Omnibus:                104.689   Durbin-Watson:          1.970
Prob(Omnibus):          0.000     Jarque-Bera (JB):        237.362
Skew:                   -1.138     Prob(JB):                2.10e-52
Kurtosis:                5.616     Cond. No.                 5.57
=====

```

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Sklearn Model Implementation

In [ ]: df\_model=df.copy()  
  
numeric\_columns = ['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA', 'Research']  
scaler = StandardScaler()  
scaler.fit(df\_model[numeric\_columns])  
df\_model[numeric\_columns] = scaler.transform(df\_model[numeric\_columns])  
  
## Min-Max Scaler Data Preparation  
# scaler = MinMaxScaler()  
# df\_model = pd.DataFrame(scaler.fit\_transform(df\_model), columns=df\_model.columns)  
# df\_model.head()

In [ ]: y = df\_model[['Chance of Admit']]  
x = df\_model.drop(['Chance of Admit'], axis=1)  
  
x\_train, x\_test, y\_train, y\_test = train\_test\_split(x, y, test\_size=0.1, random\_state=2)  
x\_train

|     | GRE Score | TOEFL Score | University Rating | SOP       | LOR       | CGPA      | Research  |
|-----|-----------|-------------|-------------------|-----------|-----------|-----------|-----------|
| 407 | -1.637030 | -1.183716   | -0.099793         | -0.882817 | 0.558125  | -1.036796 | 0.886405  |
| 291 | -1.459786 | -0.854540   | -0.975168         | -1.892906 | -1.605151 | -1.169201 | -1.128152 |
| 259 | 1.287504  | 1.943453    | 0.775582          | 1.642404  | 1.098944  | 1.263738  | 0.886405  |
| 231 | 0.224037  | -0.196189   | -0.099793         | 0.127271  | -1.064332 | -0.407873 | 0.886405  |
| 478 | 0.135415  | -0.689952   | -0.099793         | 0.632315  | 1.098944  | -0.143063 | 0.886405  |
| ... | ...       | ...         | ...               | ...       | ...       | ...       | ...       |
| 22  | 1.021637  | 1.449690    | 1.650957          | 1.642404  | 1.639763  | 1.528547  | 0.886405  |
| 72  | 0.401282  | 0.626751    | 1.650957          | 1.642404  | 1.639763  | 1.445794  | 0.886405  |
| 493 | -1.459786 | -2.006655   | -0.975168         | -0.377773 | -2.145970 | -0.589930 | 0.886405  |
| 15  | -0.219074 | -0.360777   | -0.099793         | 0.127271  | -1.064332 | -0.457525 | -1.128152 |
| 168 | -2.080142 | -1.677479   | -0.975168         | -1.387862 | 0.558125  | -1.285055 | 0.886405  |

450 rows × 7 columns

In [ ]: model = LinearRegression()  
model.fit(x\_train, y\_train)


Out[ ]: LinearRegression  
LinearRegression()

In [ ]: model.coef\_

Out[ ]: array([[0.82396356, 0.01506632, 0.00747459, 0.00072564, 0.01713408, 0.07113889, 0.01010487]])

In [ ]: imp = pd.DataFrame(list(zip(x\_test.columns, np.abs(model.coef\_[0])),  
columns=['feature', 'coeff']))  
sns.barplot(x='feature', y='coeff', data=imp)  
plt.xticks(rotation=90)

Out[ ]: ([0, 1, 2, 3, 4, 5, 6],  
[Text(0, 0, 'GRE Score'),  
Text(1, 0, 'TOEFL Score'),  
Text(2, 0, 'University Rating'),  
Text(3, 0, 'SOP'),  
Text(4, 0, 'LOR'),  
Text(5, 0, 'CGPA'),  
Text(6, 0, 'Research')])




In [ ]: model.intercept\_


Out[ ]: array([0.72157922])

In [ ]: y\_hat = model.predict(x\_test)

In [ ]: fig = plt.figure()  
x\_feature = x\_test['GRE Score']  
plt.scatter(x\_feature, y\_test, label='data')  
plt.scatter(x\_feature, y\_hat, color='orange', label='prediction')  
plt.legend()  
plt.show()



```
In [ ]: fig = plt.figure()
x_feature = x_test["University Rating"]
plt.scatter(x_feature, y_test, label='data')
plt.scatter(x_feature, y_hat, color='orange', label='prediction')
plt.legend()
plt.show()
```



```
In [ ]: ## MAE
abs(y_hat - y_test).mean()
```

```
Out[ ]: Chance of Admit      0.039666
dtype: float64
```

```
In [ ]: ## RMSE
mean_squared_error(y_test, y_hat, squared=False)
```

```
Out[ ]: 0.05877180888433925
```

```
In [ ]: ## R2 Value
r2_score(y_test, y_hat)
```

```
Out[ ]: 0.8439659610816171
```

```
In [ ]: ## Adjusted r2 value
adj_r2_func(r2_score(y_test, y_hat), x_test)
```

```
Out[ ]: 0.817960287928534
```

```
In [ ]: errors = y_hat - y_test
```

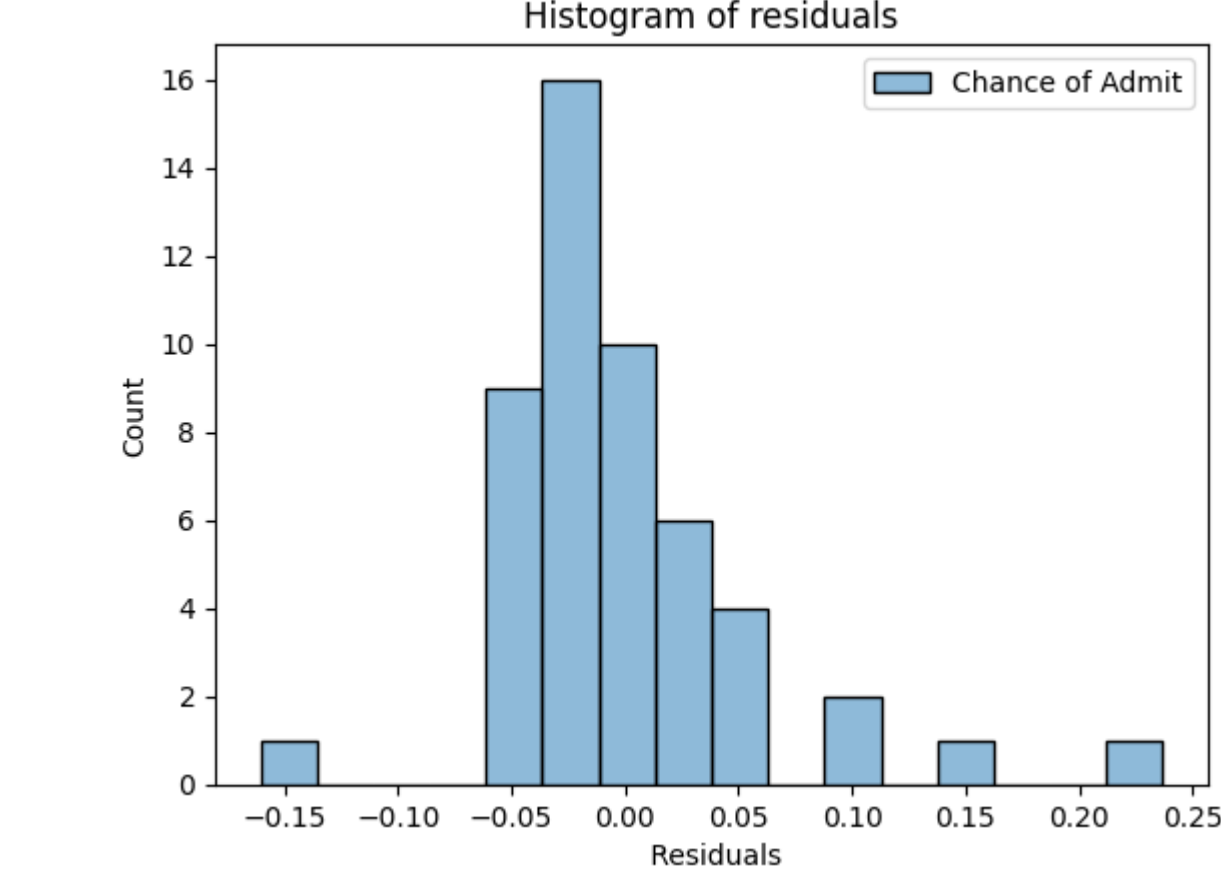
```
In [ ]: ## Means of Residuals
np.array(errors).mean()
```

```
Out[ ]: -0.0016077872077638522
```

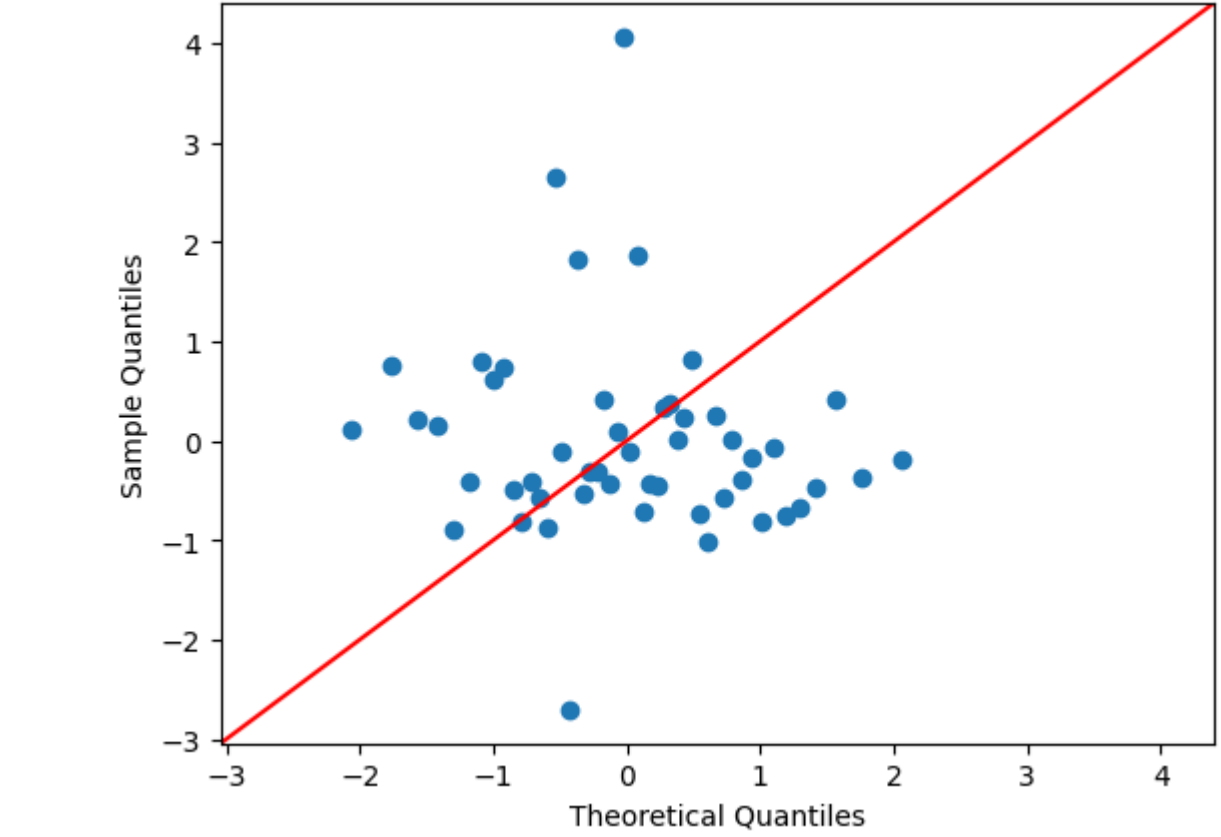
```
In [ ]: ## Normality of Residuals
```

```
sns.histplot(errors)
plt.xlabel("Residuals")
plt.title("Histogram of residuals")
```

```
Out[ ]: Text(0.5, 1.0, 'Histogram of residuals')
```




```
In [ ]: # Create a QQ plot
sm.qqplot(errors, line='45', fit=True, dist=stats.norm)
plt.show()
```



```
In [ ]: ## Normality of Residuals
res = stats.shapiro(errors)
print(res.statistic) ## Closer the value to 1, more is the normality.
0.8413031697273254
```

```
In [ ]: ## Linearity of variable / Homoscedasticity
sns.scatterplot(x=y_hat.flatten(), y=np.array(errors).flatten())
plt.xlabel("predicted Selling price")
plt.ylabel("Residuals")
plt.title("Predicted values vs Residuals")
```

```
Out[ ]: Text(0.5, 1.0, 'Predicted values vs Residuals')
```



```
In [ ]:
```